

This architecture shows a scalable microservices system designed to handle many users efficiently while maintaining reliability and performance.

First, users send requests to the system. These requests go through a Load Balancer, which distributes traffic evenly to avoid overwhelming any single service.

In the Microservices layer, different services (Microservice A, B, and C) manage specific business tasks. This separation allows each service to operate independently, making the system easier to maintain and scale.

When heavy or background tasks need processing, the microservices send them to RabbitMQ (Task Queue). RabbitMQ serves as a middle layer that temporarily stores tasks and distributes them to available workers. This keeps the system running smoothly during high traffic.

Multiple Worker Nodes (Worker 1, Worker 2, Worker N) process the tasks. These workers can scale horizontally, meaning more workers can be added when demand increases. This improves performance and keeps the system responsive.

Processed data is stored in the SQL Database, which supports read and write operations. The database connects to failover and recovery mechanisms to ensure data safety and availability.

To maintain system stability, there is a Monitoring Stack that tracks:

- System errors
- System health
- Current load

If something goes wrong, the Failover and Recovery layer manages:

- Database replication
- Worker restarts
- RabbitMQ high availability
- Backups

Overall, this architecture is built to be scalable, fault-tolerant, and reliable. It clearly separates responsibilities, distributes workload effectively, and includes monitoring and recovery systems to ensure continuous operation.